

# A comprehensive guide to Feature Selection using Wrapper methods in Python

---

 [analyticsvidhya.com/blog/2020/10/a-comprehensive-guide-to-feature-selection-using-wrapper-methods-in-python](https://analyticsvidhya.com/blog/2020/10/a-comprehensive-guide-to-feature-selection-using-wrapper-methods-in-python)

October 24, 2020

*This article was published as a part of the Data Science Blogathon.*

## Introduction

---

In today's era of Big data and IoT, we are easily loaded with rich datasets having extremely *high dimensions*. In order to perform any machine learning task or to get insights from such *high dimensional* data, **feature selection** becomes very important. Since some features may be irrelevant or less significant to the dependent variable so their unnecessary inclusion to the model leads to

- Increase in **complexity** of a model and makes it **harder** to **interpret**.
- Increase in **time complexity** for a model to get trained.
- Result in a **dumb model** with inaccurate or less reliable predictions.

Hence, it gives an indispensable need to perform *feature selection*. **Feature selection** is very crucial and must component in machine learning and data science workflows especially while dealing with *high-dimensional* datasets.

## What is Feature selection?

---

As the name suggests, it is a process of selecting the most *significant* and *relevant* features from a vast set of features in the given dataset.

For a dataset with ***d*** input features, the ***feature selection*** process results in ***k*** features such that ***k*** < ***d***, where ***k*** is the smallest set of *significant* and *relevant* features.

So ***feature selection*** helps in finding the smallest set of features which results in

- **Training** a machine learning algorithm **faster**.
- Reducing the **complexity** of a model and making it easier to **interpret**.
- Building a **sensible model** with **better prediction power**.
- **Reducing over-fitting** by selecting the right set of features.

## Feature selection methods

---

For a dataset with **d features**, if we apply the hit and trial method with all possible combinations of features then total  $(2^d - 1)$  models need to be evaluated for a *significant* set of features. It is a time-consuming approach, therefore, we use *feature selection* techniques to find out the smallest set of features more efficiently.

There are three types of *feature selection* techniques :

1. Filter methods
2. Wrapper methods
3. Embedded methods

### Difference between Filter, Wrapper, and Embedded Methods for Feature Selection

Filter methods	Wrapper methods	Embedded methods
Generic set of methods which do not incorporate a <b>specific machine learning algorithm</b> .	Evaluates on a <b>specific machine learning algorithm</b> to find optimal features.	Embeds (fix) features during <b>model building process</b> . Feature selection is done by observing each iteration of model training phase.
Much <b>faster</b> compared to Wrapper methods in terms of time complexity	<b>High computation time</b> for a dataset with many features	Sits <b>between Filter methods and Wrapper methods</b> in terms of time complexity
Less prone to <b>over-fitting</b>	High chances of <b>over-fitting</b> because it involves training of machine learning models with different combination of features	Generally used to reduce <b>over-fitting</b> by <b>penalizing</b> the coefficients of a model being too large.
Examples – Correlation, Chi-Square test, ANOVA, Information gain etc.	Examples - Forward Selection, Backward elimination, Stepwise selection etc.	Examples - LASSO, Elastic Net, Ridge Regression etc.


Filter vs. Wrapper vs. Embedded methods

In this post, we will only discuss *feature selection* using **Wrapper methods** in Python.

### Wrapper methods

In *wrapper methods*, the *feature selection* process is based on a specific machine learning algorithm that we are trying to fit on a given dataset.

It follows a *greedy search approach* by evaluating all the possible combinations of features against the *evaluation criterion*. The *evaluation criterion* is simply the performance measure which depends on the type of problem, for e.g. For *regression* evaluation criterion can be p-values, R-squared, Adjusted R-squared, similarly for *classification* the evaluation criterion can be accuracy, precision, recall, f1-score, etc. Finally, it selects the combination of features that gives the optimal results for the specified machine learning algorithm.



### Flow Chart – Wrapper methods

Most commonly used techniques under *wrapper methods* are:

1. Forward selection
2. Backward elimination
3. Bi-directional elimination(Stepwise Selection)

Too much theory so far. Now let us discuss *wrapper methods* with an example of ***the Boston house prices dataset*** available in sklearn. The dataset contains 506 observations of 14 different features. The dataset can be imported using the `load_boston()` function available in the `sklearn.datasets` module.

```
from sklearn.datasets import load_boston
boston = load_boston()
print(boston.data.shape)          # for dataset dimension
print(boston.feature_names)       # for feature names
print(boston.target)              # for target variable
print(boston.DESCR)              # for data description
```

Let's convert this raw data into a *data frame* including target variable and actual data along with feature names.

```
import pandas as pd
bos = pd.DataFrame(boston.data, columns = boston.feature_names)
bos['Price'] = boston.target
X = bos.drop("Price", 1)          # feature matrix
y = bos['Price']                  # target feature
bos.head()
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	LSTAT	Price
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	15.3	396.90	4.98	24.0
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	17.8	396.90	9.14	21.6
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0	17.8	392.83	4.03	34.7
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0	18.7	394.63	2.94	33.4
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0	18.7	396.90	5.33	36.2

Here, the target variable is Price. We will be fitting a *regression model* to predict Price by selecting optimal features through *wrapper methods*.

## 1. Forward selection

---

In *forward selection*, we start with a null model and then start fitting the model with each individual feature one at a time and select the feature with the minimum *p-value*. Now fit a model with two features by trying combinations of the earlier selected feature with all other remaining features. Again select the feature with the minimum *p-value*. Now fit a model with three features by trying combinations of two previously selected features with other remaining features. Repeat this process until we have a set of selected features with a *p-value* of individual features less than the *significance level*.

In short, the steps for ***the forward selection*** technique are as follows :

1. Choose a *significance level* (e.g. SL = 0.05 with a 95% confidence).
2. Fit all possible *simple regression models* by considering one feature at a time. Total '**n**' models are possible. Select the feature with the lowest *p-value*.
3. Fit all possible models with one extra feature added to the previously selected feature(s).
4. Again, select the feature with a minimum *p-value*. if *p\_value < significance level* then go to Step 3, otherwise terminate the process.

Now let us perform the same on Boston house price data.

```
def forward_selection(data, target, significance_level=0.05):
    initial_features = data.columns.tolist()
    best_features = []
    while (len(initial_features)>0):
        remaining_features = list(set(initial_features)-set(best_features))
        new_pval = pd.Series(index=remaining_features)
        for new_column in remaining_features:
            model = sm.OLS(target, sm.add_constant(data[best_features+
[new_column]])).fit()
            new_pval[new_column] = model.pvalues[new_column]
        min_p_value = new_pval.min()
        if(min_p_value<significance_level):
            best_features.append(new_pval.idxmin())
        else:
            break
    return best_features
```

This above function accepts data, target variable, and significance level as arguments and returns the final list of *significant features* based on *p-values* through ***forward selection***.

```
forward_selection(X,y)
```

```
#OUTPUT
['LSTAT',
 'RM',
 'PTRATIO',
 'DIS',
 'NOX',
 'CHAS',
 'B',
 'ZN',
 'CRIM',
 'RAD',
 'TAX']
```

## Implementing Forward selection using built-in functions in Python:

---

mlxtend library contains built-in implementation for most of the *wrapper methods* based feature selection techniques. SequentialFeatureSelector() function comes with various combinations of *feature selection* techniques.

```
#importing the necessary libraries
from mlxtend.feature_selection import SequentialFeatureSelector as SFS
from sklearn.linear_model import LinearRegression
# Sequential Forward Selection(sfs)
sfs = SFS(LinearRegression(),
          k_features=11,
          forward=True,
          floating=False,
          scoring = 'r2',
          cv = 0)
```

SequentialFeatureSelector() function accepts the following major arguments :

- LinearRegression() is an estimator for the entire process. Similarly, it can be any *classification* based algorithm.
- k\_features indicates the number of features to be selected. It can be any random value, but the optimal value can be found by analyzing and visualizing the scores for different numbers of features.
- forward and floating arguments for different flavors of *wrapper methods*, here, forward = True and floating = False are for *forward selection* technique.
- The scoring argument specifies the *evaluation criterion* to be used. For *regression* problems, there is only *r2 score* in default implementation. Similarly for *classification*, it can be accuracy, precision, recall, f1-score, etc.
- cv argument is for *k-fold cross-validation*.

Now let's fit the above-defined *feature selector* on the Boston house price dataset.

```
sfs.fit(X, y)
sfs.k_feature_names_      # to get the final set of features
#OUTPUT
('CRIM',
 'ZN',
 'CHAS',
 'NOX',
 'RM',
 'DIS',
 'RAD',
 'TAX',
 'PTRATIO',
 'B',
 'LSTAT')
```

## 2. Backward elimination

---

In *backward elimination*, we start with the full model (including all the independent variables) and then remove the insignificant feature with the highest *p-value* (> *significance level*). This process repeats again and again until we have the final set of *significant* features.

In short, the steps involved in *backward elimination* are as follows:

1. Choose a *significance level* (e.g. SL = 0.05 with a 95% confidence).
2. Fit a full model including all the features.
3. Consider the feature with the highest *p-value*. If the *p-value* > *significance level* then go to Step 4, otherwise terminate the process.
4. Remove the feature which is under consideration.
5. Fit a model without this *feature*. Repeat the entire process from Step 3.

Now let us perform the same on Boston house price data.

```
def backward_elimination(data, target, significance_level = 0.05):
    features = data.columns.tolist()
    while(len(features)>0):
        features_with_constant = sm.add_constant(data[features])
        p_values = sm.OLS(target, features_with_constant).fit().pvalues[1:]
        max_p_value = p_values.max()
        if(max_p_value >= significance_level):
            excluded_feature = p_values.idxmax()
            features.remove(excluded_feature)
        else:
            break
    return features
```

This above function returns the final list of *significant features* based on *p-values* through ***backward elimination***.

```
# OUTPUT
['CRIM',
 'ZN',
 'CHAS',
 'NOX',
 'RM',
 'DIS',
 'RAD',
 'TAX',
 'PTRATIO',
 'B',
 'LSTAT']
```

## Implementing Backward elimination using built-in functions in Python:

---

The same SequentialFeatureSelector()function can be used to perform *backward elimination* by disabling the forward argument.

```
#Sequential backward selection(sbs)
sbs = SFS(LinearRegression(),
          k_features=11,
          forward=False,
          floating=False,
          cv=0)
sbs.fit(X, y)
sbs.k_feature_names_
# OUTPUT
('CRIM',
 'ZN',
 'CHAS',
 'NOX',
 'RM',
 'DIS',
 'RAD',
 'TAX',
 'PTRATIO',
 'B',
 'LSTAT')
```

## Additional Note

---

Here we are directly using the optimal value of k\_features argument in both *forward selection* and *backward elimination*. In order to find out the optimal number of *significant* features, we can use *the hit and trial* method for different values of k\_features and make the final decision by plotting it against the model performance.


```
sfs1 = SFS(LinearRegression(),
            k_features=(3,11),
            forward=True,
            floating=False,
            cv=0)
sfs1.fit(X, y)
```

The same visualization can be achieved through plot\_sequential\_feature\_selection()function available in mlxtend.plotting module.

```

from mlxtend.plotting import plot_sequential_feature_selection as plot_sfs
import matplotlib.pyplot as plt
fig1 = plot_sfs(sfs1.get_metric_dict(), kind='std_dev')
plt.title('Sequential Forward Selection (w. StdErr)')
plt.grid()
plt.show()

```



Here, on the y-axis, the performance label indicates the *R-squared* values for the different numbers of features.

### 3. Bi-directional elimination(Step-wise Selection)

---

It is similar to *forward selection* but the difference is while adding a new feature it also checks the *significance* of already added features and if it finds any of the already selected features insignificant then it simply removes that particular feature through *backward elimination*.

Hence, It is a combination of *forward selection* and *backward elimination*.

In short, the steps involved in *bi-directional elimination* are as follows:

1. Choose a *significance level* to enter and exit the model (e.g.  $SL_{in} = 0.05$  and  $SL_{out} = 0.05$  with 95% confidence).
2. Perform the next step of *forward selection* (newly added feature must have *p-value*  $< SL_{in}$  to enter).
3. Perform all steps of *backward elimination* (any previously added feature with *p-value*  $> SL_{out}$  is ready to exit the model).
4. Repeat steps 2 and 3 until we get a final *optimal* set of features.

Let us perform the same on Boston house price data.



```

def stepwise_selection(data, target, SL_in=0.05, SL_out = 0.05):
    initial_features = data.columns.tolist()
    best_features = []
    while (len(initial_features)>0):
        remaining_features = list(set(initial_features)-set(best_features))
        new_pval = pd.Series(index=remaining_features)
        for new_column in remaining_features:
            model = sm.OLS(target, sm.add_constant(data[best_features+
[new_column]])).fit()
            new_pval[new_column] = model.pvalues[new_column]
        min_p_value = new_pval.min()
        if(min_p_value<SL_in):
            best_features.append(new_pval.idxmin())
            while(len(best_features)>0):
                best_features_with_constant = sm.add_constant(data[best_features])
                p_values = sm.OLS(target,
best_features_with_constant).fit().pvalues[1:]
                max_p_value = p_values.max()
                if(max_p_value >= SL_out):
                    excluded_feature = p_values.idxmax()
                    best_features.remove(excluded_feature)
                else:
                    break
            else:
                break
        return best_features

```

This above function returns the final list of *significant* features based on *p-values* through *bi-directional elimination*.

```
stepwise_selection(X,y)
```

```

# OUTPUT
['LSTAT',
'RM',
'PTRATIO',
'DIS',
'NOX',
'CHAS',
'B',
'ZN',
'CRIM',
'RAD',
'TAX']

```

## Implementing bi-directional elimination using built-in functions in Python:

---

The same SequentialFeatureSelector() function can be used to perform *backward elimination* by enabling forward and floating arguments.

```
# Sequential Forward Floating Selection(sffs)
sffs = SFS(LinearRegression(),
          k_features=(3,11),
          forward=True,
          floating=True,
          cv=0)
sffs.fit(X, y)
sffs.k_feature_names_
# OUTPUT
('CRIM',
 'ZN',
 'CHAS',
 'NOX',
 'RM',
 'DIS',
 'RAD',
 'TAX',
 'PTRATIO',
 'B',
 'LSTAT')
```

## End Notes

---

In this article, we saw different kinds of Wrapper methods for feature selection with implementation using mlxtend library in Python.